


System virtualization

Operating system (OS) and its Goals

- **Resource Management**
 - OS controls how processes share hardware (CPU, memory, storage, network, etc.)
- **Abstraction**
 - Hide underline details
 - Provide usable interfaces
- **Protection and Privacy**
 - Process cannot access other process data



资源管理，向下管理各种硬件资源

向上抽象，让上层以为自己独占资源，可以正常使用资源。

保护、隐私。不能接触到其他进程的数据。

操作系统与硬件之间的接口是 ISA。


操作系统与进程之间的接口是 ABI（系统调用）。

进程抽象

1. 每个进程有自己的私有地址空间
2. 操作系统调度进程占用 CPU（分时）
3. 进程通过系统调用来使用 OS 的服务

Memory Management Unit (MMU)

- OS configures page tables for processes
- Install a page table by modifying CR3 register
- MMU checks a page table when translating virtual addresses



模式转换的方式：

Three Types of Mode Switches

- **System Calls**
 - Process requests a system service, e.g., exit
 - Like a function call, but outside the process
 - Do not have the address of the system function to call
- **Exceptions**
 - Internal synchronous event in process triggers context switch
 - Protection violation (segmentation fault), Divide by zero, ...
- **Interrupts**
 - External asynchronous event triggers context switch
 - Timer, I/O device
 - Independent of user process

OS 提供的系统隔离

进程之间的隔离

技术：page tables, context switch, file abstraction


进程和内核之间的隔离

CPU 特权级

VMM

Virtual Machine Monitor

- A VMM (aka Hypervisor) provides a system virtual machine to each OS
- VMM can run directly on hardware (type-1) or on another OS (type 2)
- **Hardware virtualization**
 - VT-x: root and non-root mode
 - EPT
 - IOMMU and SR-IOV



两种 VMM

Type1: 直接运行在硬件上


Type2: 运行在已有的操作系统上

Type2 可以只关注虚拟化的部分，其他例如设备驱动的功能可以复用 OS。

和 OS 对进程的抽象类似，VM 抽象：

Virtual Machine (VM) Abstraction


- Each VM has a private address space
 - Provided by the hypervisor
 - Cannot access other VM spaces
- The hypervisor schedules processes into cores
 - Each process has a scheduling time slice
- A VM can invoke hypervisor services via hypercalls



System ISA: 可以访问敏感资源的 ISA

System ISA

- **Access Sensitive Registers**
 - CR0, CR3, CR4...
- **Control CPU**
 - Example: HLT
- **Control virt/phy memory**
 - Configure & Install PT
- **Control Devices**
 - DMA, Interrupts




CR3: 进程对应的页表地址。

系统虚拟化的步骤：

Procedure of System Virtualization

- **Step 1**
 - Catch all system ISA of a VM
- **Step 2**
 - Provide three functions
 - Virtualize CPU
 - Virtualize Memory
 - Virtualize I/O
- **Step 3**
 - Resume VM



首先要捕获 VM 的所有的系统 ISA 识别并提供相关的服务。回到虚拟机中继续执行。

CPU 虚拟化

如果简单把 OS 当作普通的 Application, 在 OS 上去跑, 问题是客户机 OS 要执行特权指令, 但是其所在特权级不能执行特权指令。

Solution: Trap & Emulate

- **Trap:** running privilege instructions will trap to the VMM
- **Emulate:** those instructions are implemented as functions in the VMM

问题是, 架构并非是“严格可虚拟化”的。

Problems of Trap & Emulate

- Not all architectures are “strictly virtualizable”
- An ISA is strictly virtualizable if, when executed in a lesser privileged mode:
 - All instructions that access privileged state trap
 - All instructions either trap or execute identically

X86 结构中, 一些特权级指令, 在用户特权级下运行时, 不会正常生效, 只是静默地跳过去, 什么都不发生, 不会 trap。

解决方法:

How to Deal with the 17 Instructions?

1. **Instruction Interpretation:** emulate them by software
2. **Binary translation:** translate them to other instructions
3. **Para-virtualization:** replace them in the source code
4. **New hardware:** change the CPU to fix the behavior

Sol-1: 指令翻译

Sol-1: Instruction Interpretation

- **Emulate Fetch/Decode/Execute pipeline in software**
 - Emulate all the system status using memory
 - E.g., using an array `GPR[8]` for general purpose registers
 - None guest instruction executes directly on hardware
- E.g., Bochs

把每一条指令都用软件模拟执行。最大的问题就是慢。

Sol-2 二进制翻译

Sol-2: Binary Translator

- **Translate before execution**
 - Translation unit is basic block (why?)
 - Each basic block -> code cache
 - Translate the 17 instructions to function calls
 - Implemented by the VMM
- E.g., VMware, Qemu

在执行之前提前翻译好。

翻译的单元是基础块。

把 17 个指令翻译成 call

Sol-3 半虚拟化

Sol-3: Para-virtualization

- **Modify OS and let it cooperate with the VMM**
 - Change sensitive instructions to calls to the VMM
 - Also known as hypercall
 - Hypercall can be seen as trap
- E.g., Xen
 - Was widely used by industry like Amazon's EC2

改操作系统, 把敏感指令换成 hypercall。

(具有可行性是因为要改的敏感指令少, 常用的操作系统少)

Sol-4 硬件虚拟化

把 CPU 特权级进一步细化。


在 4 个特权级之外的维度, 加了 root 和 non-root 的模式。

Sol-4: Hardware Supported CPU Virtualization

- **VMX root operation:**
 - Full privileged, intended for Virtual Machine Monitor
- **VMX non-root operation:**
 - Not fully privileged, intended for guest software

Both forms of operation support all four privilege levels from 0 to 3

Sol-4: Hardware Supported CPU Virtualization



VMCS: VM control structure

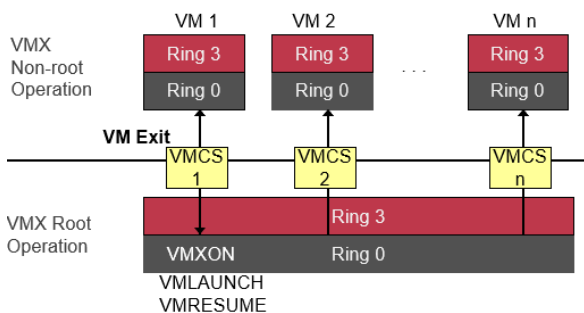
VT-x VMCS

The VMCS consists of six logical groups:

- **Guest-state area:** Processor state saved into the guest-state area on VM exits and loaded on VM entries.
- **Host-state area:** Processor state loaded from the host-state area on VM exits.
- **VM-execution control fields:** Fields controlling processor operation in VMX non-root operation.
- **VM-exit control fields:** Fields that control VM exits.
- **VM-entry control fields:** Fields that control VM entries.
- **VM-exit information fields:** Read-only fields to receive information on VM exits describing the cause and the nature of the VM exit.

VT-x 工作流程

VT-x Workflow



首先执行 VMXON 指令，告诉 CPU 要使用虚拟化了。执行 VMLAUNCH，就会进入虚拟机中执行。遇到特权指令后会 VM Exit，回到 VMM，执行完之后执行 VMRESUME，再回到 VM。每个虚拟机有一个 VMCS，告诉有哪些指令需要下陷。

内存虚拟化

扩展了内存地址种类：

Terminology: 3 types of address now

- **GVA->GPA->HPA** (Guest virtual. Guest physical. Host physical)
- Guest VM's page table contains GPA

解决方法：

- Traditional solutions: **shadow paging & direct pa**
- Today's solution: **new hardware**

扩展页表


Hardware implementation

- Intel's EPT (Extended Page Table)
- AMD's NPT (Nested Page Table)

Another table

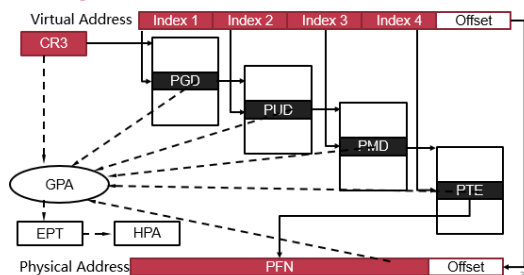
- EPT for translation from **GPA to HPA**
- EPT is controlled by the hypervisor
- EPT is per-VM

EPT 用来把 GPA 翻译成 HPA



原本的地址翻译过程：

Any GPA is translated to HPA



5 次内存访问。

EPT 翻译与此类似。

意味着需要 20 次 (4*5) 内存访问。

IO 虚拟化

I/O Virtualization

Goal

- Multiplexing device to guest VMs

Challenges

- Each guest OS has its own driver
- How can one device be controlled by multiple drivers?
- What if one guest OS tries to format its disk?

四个方案：

Solutions for I/O Virtualization

1. **Direct access:** VM owns a device exclusively
2. **Device emulation:** VMM emulates device in software
3. **Para-virtualized:** split the drivers to guest and host
4. **Hardware assisted:** self-virtualization device

Sol-1 直通独占

问题：虚拟机可以通过 DMA 让设备访问到其他虚拟机的内存

解决方法：设备页表，IOMMU

IOMMU: Page Tables for Devices

- **Allow guest OS direct access to underlying device**
 - Guest just reuses its own device driver
- **Q: What if a VM asks device to access memory of other VMs?**
 - It is possible because device accesses HPA in DMA
 - Thus a device can access any memory
- **Solution: Page tables for devices**
 - Another MMU: IOMMU for devices
 - Q: what addresses will IOMMU translate?

优点：快，简化 VMM

缺点：guest 直接操作硬件接口（难以做 VM 迁移）（不走 VMM）。需要更多设备。

Direct Access Device Virtualization

- **Positives**
 - Fast, since the VM uses device just as native machine
 - Simplify monitor: limited device drivers needed
- **Negatives**
 - Hardware interface visible to guest (bad for migration)
 - Interposition is hard by definition (no way to trap & emulate)
 - Now you need much more devices! (image 100 VMs)

Sol-2 模拟设备

用软件模拟设备逻辑


Emulated Devices

- **Positives**
 - Platform stability (good for migration)
 - Allows interposition
 - No special hardware support is needed
- **Negatives**
 - Can be slow (it's software emulated)

Sol-3 半虚拟化

Sol-3: Para-Virtualized Devices

- **VMM offers new types of device**
 - The guest OS will run a new driver (front-end driver)
 - The VMM will run a back-end driver for each front-end
 - The VMM will finally run device driver to drive the device



两边互相知道，可以更好的配合，比如数据可以在前端 batch 之后再给后端。

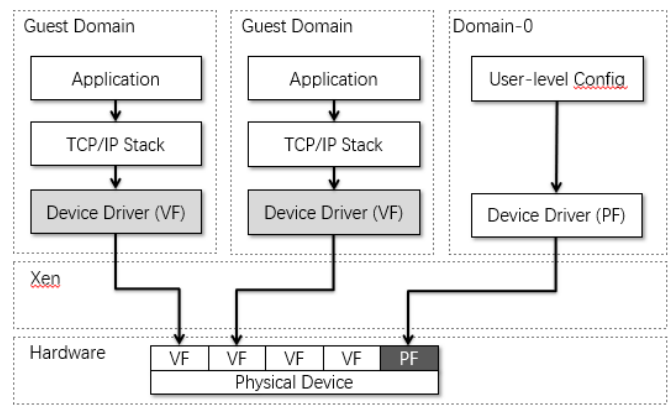
因此快。

Sol-4 硬件虚拟化

Sol-4: Hardware Support for I/O Virtualization

- **VMM**
 - An SR-IOV-capable device can be configured to appear in the PCI configuration space as multiple functions
- **VM**
 - The VMM assigns one or more VFs to a VM by mapping the actual configuration space of the VFs to the configuration space presented to the virtual machine by the VMM

SR-IOV：硬件的一种虚拟化标准



PF 具有管理功能，被 VMM 控制。

总结表格：

Virtualization Technologies

Virtualization	Software Solution	Hardware Solution
CPU	<ul style="list-style-type: none"> Trap & Emulate Instruction interpretation Binary translation Para-virtualization: Replace 17 instructions 	<ul style="list-style-type: none"> VT-x <ul style="list-style-type: none"> Root / non-root mode VMCS
Memory	<ul style="list-style-type: none"> Shadow page table Separating page tables for U/K Para-virtualization: Direct paging 	<ul style="list-style-type: none"> EPT
Device	<ul style="list-style-type: none"> Direct I/O Device emulation Para-virtualization: Front-end & back-end driver (e.g., virtio) 	<ul style="list-style-type: none"> IOMMU SR-IOV

52

Topic: Isolation Via VMM

SeCage [CCS '15]: Isolate Sensitive Code

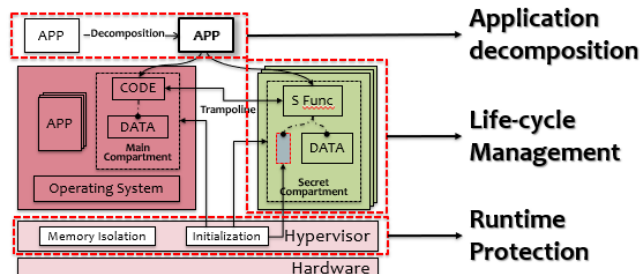
Problem

Private data leakage caused by memory vulnerabilities. Example: **HeartBleed** ❤️

Solution

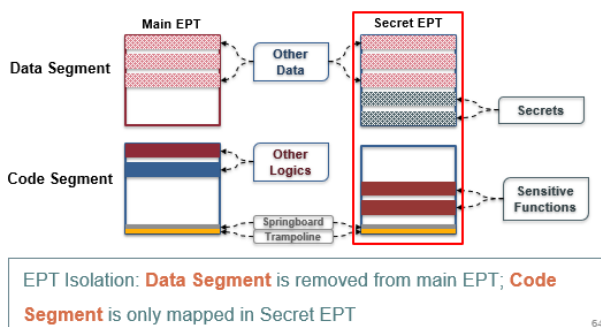
- Isolation Domain: Put sensitive code via an isolated VM
- Support different applications, including **OpenSSL**
- Utilize hardware features to switch between the application and its isolation domain, causing only 8% overhead

SeCage [CCS '15]: Isolate Sensitive Code



秘密数据部分摘出来，为其单独分配一个 EPT，使得 OS 不能直接访问。

SeCage: Different EPTs for two Parts



缺点：如果秘密数据很多，被频繁用，那么每次访问需要到 VMM 里换页表，很慢。


解决方案：intel 加了一个指令 VMFUNC，可以在 VM 里

换页表，不需要 VM Exit

CloudVisor：在恶意的 hypervisor 情况下保护 VM

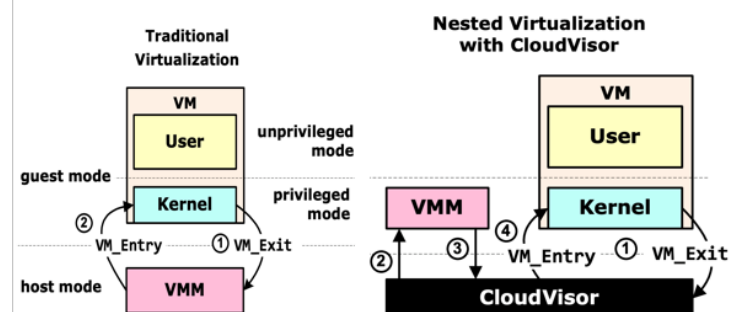
Key Idea

- Separating security protection from VM hosting
- Add another layer of indirection
 - Nested virtualization




管理和安全分离。VMM 被降权，CloudVisor 在最底层，起保护作用。

General Workflow



开销过大，慢

The Cost of Protection: Excessive VM Exits




Operation	Times
Hypercall	$\geq 2X$
EPT Violation	$2 - 6 X$
DMA Operation	$\geq 2X$

改进 CloudVisor-D

Architecture of CloudVisor-D

- A tiny nested hypervisor in root mode
- A Guardian-VM for each VM in non-root mode
- Most VM ops offloaded to Guardian-VM
 - Hypercalls
 - Memory virtualization
 - I/O operations



Guardian-VM 在 non-root mode, 承担保护安全的作用。用 VMFUNC 切换页表。

数据流控制流分离

Language Virtual Machines

Native 语言，比如 C++，编译、部署。

便利性问题：在多个平台下，多个 ISA，需要每个平台都有对应的编译器，编译出对应的 binary。部署环境越来越复杂，不同的操作系统，不同的硬件平台。

安全性问题：


- Out-of-bound access

```
char * chs = malloc(10);
Char ch = chs[10]; // oops!
```

- Use-after-free

```
free(chs);
Char ch = chs[5]; // oops!
```

解决思路：加一层抽象



一次编译，处处执行。

安全问题的解决：

Security Enhancement

- Out-of-bound access ← runtime check

```
char * chs = malloc(10);
Char ch = chs[10]; // ArrayOutOfBoundsException
```

- Use-after-free ← no free at all!

```
free(chs);
Char ch = chs[5]; // No problem!
```


托管给 JVM 运行，JVM 负责运行时检查、内存管理的工作。

执行模式：解释执行

在 Java 栈上执行。

Stack-based Interpretation


- All instructions (bytecode) are operated on the stack



问题：efficiency

Problem: Efficiency

- Too many memory operations (≥ 1 per bytecode)
 - Low utilization of registers
- Missing optimization opportunities




内存操作太多，对寄存器的利用率低，难以优化。

解决方法：JIT compiler

Solution: JIT (Just-In-Time) Compiler


- Hot code will be compiled and executed
 - "Hot" means that it has been executed for times




并不是所有的代码都编译，有一个 profiler 进行统计，偶尔执行的就直接解释执行，经常执行的就进行编译。

JIT：激进 or 保守


- Choice 1: only compile A



- Choice 2: inline B




• Choice 3: inline all



在多态的情况下


• Choice 1: compile A only

- Suitable for all cases where all methods are hot



• Choice 2: compile only for one path

- Suitable for cases where only one path is hot



激进和保守的权衡：

代码执行效率和运行时编译开销的权衡

JIT: Aggressive Or Conservative?

• Summary: A tradeoff between code efficiency and runtime overhead



内存管理

数据放在 Java Heap 里

Free Memory Management

GC：识别死对象；管理堆空间。


堆的组织-1: free list

Free 的块组织成链表

在 Java 中不受欢迎

• Not welcomed in Java

- Fragmentations
- Multi-threaded contention



堆的组织-2: 连续空间

Heap Layout 2: Contiguous Space

• Free space is always contiguous

- A bump pointer to mark how much has been used
- Allocation: Lock-free atomic instructions (CAS)




为了维持连续空间，需要在垃圾回收之后把 live object 压在一起。

小优化：local heap

每个线程有一个本地堆，省掉 CAS 开销。

• Even faster allocation: local heap

- Allocate a large portion with CAS
- Then allocate locally



GC 算法-1: 引用计数

记录有多少指向该对象的引用。引用计数变为 0 时就行回收。

在 JAVA 中不受欢迎的原因：性能差。和 free list 有较强耦合。不能处理循环引用。

GC 算法-2: tracing 追踪

通过图遍历的方法把所有的活对象都找到。

从“root”开始遍历。

Root 包含：线程栈上的对象。全局对象，包含每个类的静态变量等。

可以解决循环引用问题。

Tracing GC 应用比较广泛。

Tracing GC is Popular

- Can be integrated with different layout
 - Free-list: Mark-Sweep (Boehm GC)
 - Contiguous: Mark-Copy (PS, G1, Shenandoah...)
- Can also be used for different purposes
 - Throughput-oriented: Stop-the-world tracing
 - Latency-oriented: concurrent tracing

吞吐量优先 or 延迟优先 ?


Throughput-oriented GC

- Design goal: high GC throughput
 - Stop-the-world: Java threads must be paused during GC
 - Task-based parallelism: dividing collections into tasks
- Consider latency together with throughput
 - Generational
- 在 GC 期间暂停线程的运行，以提高吞吐量。
- 把垃圾回收拆分成小的任务，并行处理。
- 分代。

Stop the world:

JVM leverages safepoint to pause all Java threads

- Java threads queue up for a lock held by GC threads



Aiming at high GC throughput


- CPUs are monopolized by GC threads
- No coordination between GC and Java threads

GC 线程可以充分利用 CPU;

无需考虑 GC 对 Java 线程的影响。

基于任务的并行

遍历过程：可以通过划分 roots 来分成多个任务，多个 GC 线程并行执行。



问题：根据起点分配，可能会造成负载不均衡。
可以通过 work-stealing 实现动态平衡。

考虑时延:

GC 暂停的时间和 live object 的大小正相关。当 live object 特别大时暂停的时间会很长。

因此不能等到完全没有空间了才回收。

Solution: 分代

分为新生代、老生代。

刚分配的对象在新生代，存活时间长的对象在老生代。


Solution: Generational GC

• Dividing Java heap into multiple spaces

- In PS it has two spaces: *young* and *old*
- The size of young gen can be small and fixed

• GC is also two-fold

- Minor GC: only collecting young-space (fast and frequent)
- Major GC: collecting both spaces (slow but infrequent)



GC 也分为两个过程，对新生代区域的 GC 和对全局的 GC。

基于的假设：大多数对象死的很快。因此对新生代区域的 GC 效率很高。

但是，有些场景下“分代假设”不成立，比如大数据场景 spark Hadoop。对象可能都移到老生代里去了。

分代的问题：跨区的引用

我们只想对新生代区进行遍历、回收。但是有些引用是从老生代指向新生代的。

方法是把新生代中被老生代引用的对象标为 root。


那么如何识别有哪些是跨区的引用呢？

- 扫描整个老生代？（那还分个锤子代）
- 记住所有的引用（开销太大）

方法：card table

Solution: Card Table

- **Dividing old-space into many *regions***
 - Using 1 bit (card) in a table for each region
- **When a cross-space write happens, dirty the card**
 - Scanning dirty cards only during minor GC



是一个折中的办法。

面向时延的 GC

目标是更短的停顿，更小的延迟

Design goal: shorter pauses/app latency


- Concurrent marking
 - Choose which parts are valuable for collection
- Concurrent collection
 - Collecting when application threads are active

并行标记

Java 线程和标记过程并行，处理并发更新的方法是 Java 线程会将更新内容写入 logs，标记线程会定期消费日志。

Handling concurrent update

- Java threads will write update into *logs*
- Marking threads will consume the logs periodically




并发标记的好处：

- 标记不需要暂停 Java 线程
- 标记的结果可以指导 GC 线程进行更高效的垃圾回收。

An Example: Region-based Profiling

- **The whole heap is split into many *regions***
- **GC threads estimate per-region collection cost according to concurrent marking**
 - Meeting the “soft limits” from users



可以根据标记的结果预估每个区域垃圾回收的时间。以便满足用户对时延的要求。

并行回收

会使得暂停时间非常短。


要解决的主要问题：双拷贝，引用更新。

双拷贝问题：GC 时要进行拷贝，存在一个时间窗口使得两个拷贝同时存在，可能会使得不同的线程指向的是不同的拷贝，此时如果发生更新，会不一致。

解决：间接指针

• Solution: indirect pointer

- Always point to the newest version
- All read/writes will be forwarded to the newest one




解决：双写，要解决双写的原子性问题。

引用更新问题：

In-Reference Updates

- **Only part of heap space is collected each time**
- **Only out-references are updated**



In-Reference Updates

- **Design 1: Stop-The-World update**
 - Large pause time (violating the goal of Shenandoah)
- **Design 2: memorizing all in-references**
 - Large memory overhead (all references are duplicated)
- **Design 3: lazily updated in next phases**
 - A next marking phase will scan the whole heap

较好的是 lazily update，在下一轮扫全栈的时候才真正回收，在此之前用间接指针占位。

总结：


Summary: Throughput vs. Latency

- **PSGC: Stop-The-World**
 - Throughput-oriented
 - Large GC pauses
- **G1GC: adjustable & partially-concurrent**
 - Concurrent marking
 - Controllable GC pauses
- **Shenandoah: mostly-concurrent**
 - Concurrent marking & collection
 - Ultra-low GC pauses
 - Hurting application throughput



Core Concept: Java Memory Model (JMM)

- Dictates which writes a read can observe
- Specifies 'happens-before' partial order
 - Accesses in single thread are totally ordered
 - Accesses in different threads are ordered by locks




94

引入 JVM 这一层抽象的好处是可以捕捉更多的应用相关的语义。坏处是运行时的行为更加复杂（比如 GC）。基于这些好处和坏处，有以下一些工作。
12 年的工作：COMET

The Goal of COMET

- **Transparent offloading**
 - No programming effort, arbitrary apps
- **Fine-grained offloading**
 - Restricting transferred bytes
- **Robust offloading**
 - Resisting network failures

COMET builds a distributed shared memory (DSM) between phones and clouds



传统 DSM 是基于页的内存管理。
在 JVM 下的好处是可以做**细粒度**的 DSM。


单线程内部的读写操作是按顺序的。
多个线程之间的顺序是锁来保证

- Used to establish 'happens-before' relation
- Synchronizes
 - Bytecode sources
 - Java thread context (all frames of all threads, including pc/registers/method)
 - Java heap (only dirty fields in tracked set)

关于 VM 带来的问题的工作：Semeru

Resource Disaggregation


- **Traditional data center: monolithic server model**
 - Each server hosts all types of hardware resources
 - Inducing resource under-utilization



可能因为某一种资源用满了造成其他资源的浪费。

Resource Disaggregation

- **Disaggregated data center: the next generation?**
 - Each server(s) contains different types of resources
 - Better resource utilization




1

一种简单的模型，CPU server + memory server

A simplified model: CPU server + memory server

- Applications are running on CPU server
- Data are stored on memory servers
- CPU servers keep a local memory cache (so locality matters)




105

Conclusion

- Language virtual machines: an "indirection" for portability and safety
- A typical example: JVM
 - Code execution: interpreter & JIT compiler
 - Memory management (GC): basics & modern GC designs
 - JVM in systems: COMET & Semeru


发现资源分离对 JVM 不友好。

- Workload: KMeans in Spark
 - One CPU server, two memory servers, G1GC
- Result: 18.9X for 25% cached memory
 - Full GC: 126.1X
 - Reason: too many data fetching




Semeru Overview

- A unified Java heap among all servers
 - Each memory server runs a LJVM



Key Insight: Close-Data GC

- Let memory servers handle most GC tasks
- Now GC is divided into two sub-phases:
 - Memory server concurrent tracing (MSCT)
 - CPU server STW collection (CSSC)



108

思路：memory server 处理大多数 GC 工作
CPU server 把 root 发到 memory server