

07 SQL

- DBMS 和 KV 数据库比较: In a key-value store, all the relationships and operations are managed **by the programmer** (p29左右)

关系模型 (p32起)

关系relation, 元组tuple

怎么定位一个tuple? Table name + primary key

Short summary of relational model

Goal: hide the implementation detail of data behind a clean interface

- Yet, the interface is expressive enough
- The interface is also pr

Relational model

- Data model is based on set
- Data query language is based on set algebra

The query language (e.g., SQL) is **declarative**

- The developers only specify what they want

DML (Data Manipulation Languages)

- procedural 过程式、declarative 声明式

关系代数 (p40-50)

复杂数据关系的表示 (p59起)

- Duplication and normalizing 重复和规范化 (p65)
- representing one-to-many 表示一对多关系 (p66)
 1. 解决方案1: 添加新表和外键
 2. 解决方案2: **document model 文档模型** (一对多关系)

Document model 文档模型 (e.g. JSON)

- schema flexibility: 任意修改schema
 - 关系数据库中的schema更改很慢

关系模型和文档模型对比

Summary: relational model vs. document model

Document model is a representative model in NoSQL databases

Benefits of **document model**

- Better locality
- Schema flexibility

Benefits of **relational model**

- Join supports
- Better modeling many-to-one & many-to-many relationships

OldSQL → NoSQL → NewSQL → HTAP

OldSQL (p81)

- **OldSQL = Relational Model + SQL + ACID**
- OLTP transaction
 - 生命周期短
 - 需要的数据量小

OldSQL = Relational Model + SQL + ACID

The architectural or historical baggage of SQL

Abstraction

- Relational + Transaction

Semantic

- ACID

Architecture

- Mostly focus on Single-node & On-disk

NoSQL (p91)

Scale Horizontally with Middleware 使用中间件横向扩展

- [√] Read/Write single Record
- [×] No Distributed Transaction & Join
 - Unavailable when: Changing schema - Server failover
 - Data Scaling & Re-sharding

tradeoff (p95-99)

- 简化数据模型 => 降低复杂度
- 弱化transaction, 异步复制 => 牺牲一致性, 换取scalability和availability

NoSQL - Build from scratch

How does NoSQL make trade off?

- Specific (simplified) data model
 - Weaken Transaction
 - Async Replication
- ➔ Reducing the **complexity** of Relational Model (Mainly due to Join)

↓ ↓
Sacrifice consistency for **scalability & availability**

NewSQL (p105)

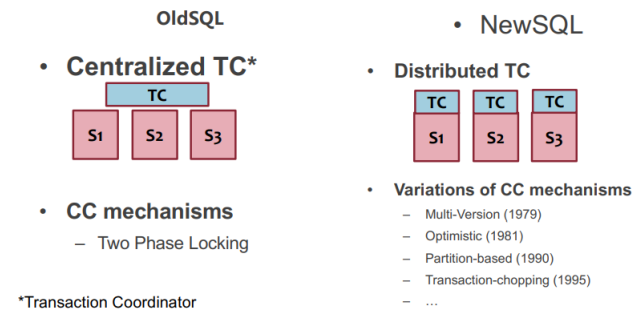
定义: A class of modern RDBMS, which provides:

- NoSQL's Scalability
- SQL's ACID Transaction & Relational Model

怎么扩展scale?

- **Shared-nothing Partitioning**

Concurrency Control



HTAP(Hybrid Transactional/Analytical Processing) (p131)

- **real-time analytics on fresh data**
- OLAP (ML)

L1 Lightning (p137-140)


L1 Lightning

A service providing Real-Time Analytic upon

- Multiple unmodified transactional database
 - Google F1 and Spanner
- Transparent to clients
 - Only need to specify which table F1 Lightning targets

Customers cannot easily migrate to new HTAP Systems

Architecture of L1 Lightning



08 Transactions

- **Atomicity:** TX is either performed entirely or not performed at all. (**rollback回滚**)
- **Consistency:** Transaction must change the data from a consistent state to another
- **Isolation:** Two concurrently executed transactions are isolated from each other
- **Durability:** Once a transaction is committed, its changes must durably stored to a persistent storage

如何保证? (p28)

- I: 一致性控制方法

Serializability (p31)

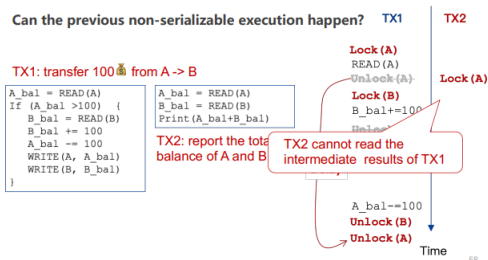
- 是理想化的：并发事务T1,T2,...,TN好像是按顺序执行的一样
- 如何检查是不是serializability? replay the concurrent execution to a serial of reads/writes (例子p35-42)

Serializability的实现 (p45)

使用锁lock

- 全局锁global lock: 一次只能一个TX, 不并发 (性能差)
- 简单细粒度锁simple fine-grained lock: 每个record有一个锁 (锁容易放太早, 不正确)
- **Two-phase locking 2PL**: TX commit之后再放所有的锁 -> **保证Serializability**

Solution#3: Two-phase locking



2PL的问题——deadlock (p61)

解决方法:

Resolving deadlock

1. **Acquire locks in a pre-defined order**
 - Not support general TX: TX must know the read/write sets before execution
2. **Detect deadlock by calculating the conflict graph**
 - If there is a cycle, then there must be a deadlock
 - Abort one TX to break the cycle
 - High cost for detection
3. **Using heuristics (e.g., timestamp) to pre-abort the TXs**
 - May have false positive, or live locks

Optimistic concurrency control -- OCC (p67)

- Executing TXs **optimistically** w/o acquiring the lock
- Checks the **results of TX** before it commits
 - If violate serializability, then **aborts & retries**

OCC Executes a Transaction in 3 Phases

Phase 1: Concurrent local processing

- Reads data into a read set
- Buffers writes into a write set

Phase 2: Validation in critical section

- Validates whether serializability is guaranteed:
- Has any data in the read set been modified?

Phase 3: Commit the results in critical section or abort

- Aborts: aborts the transaction if validation fails
- Commits: installs the write set and commits the transaction

具体例子 (p69-80)

好处:

- phase1: Operates in **private** workspace; **rare inter-thread** synchronization (optimistic)
- phase2, 3: Needs synchronization, but usually very **short at low contention**

问题:

- False Aborts 错误的abort
- livelock: high contention情况下, 一直abort, 没有progress

Summary of realizing serializability

Pessimistic methods

- Presume that interference is likely
- Prevent any possibility of conflict actively
- E.g., global lock, 2-phase locking

Optimistic methods

- Allow write in any order and at any time
- If detect conflict, then "sorry, conflict write, please abort, clear the history and then retry"
- E.g., OCC

Modern Transaction Systems (p89)

HTM (p92)

intel RTM (p96)

Fun facts about RTM

How does Intel implement RTM?

- Basically, OCC!
- Use CPU cache to track TX' s read/write sets

Why efficient?

- Cache is a perfect place for tracking TX' s temporal updates
- Leverage existing cache coherence protocol to detect conflicts: reuse existing multi-core hardware to implement transactional memory!

Hardware support for transactional memory总结

- Easy programming model for the programmer
- Good performance if using properly
- However, the programmer should handle its pitfalls

DBX (p111)

a TX system to use **RTM for acceleration**, but **avoids its pitfalls** for TXs

Conclusion of DBX

RTM is a promising feature provided by Intel CPU

RTM alone is not sufficient for TXs

DBX provides a study of how to use RTM to accelerate in-memory databases

RTM can simplify the system building and get comparable performance

- Limitations of RTM force us to craft the transaction region and memory access pattern carefully

ROCOCO (p135)

回顾：优化TX的方向？ Improve the TX algorithms properties

- E.g., better deadlock detection algorithms in 2PL
- E.g., reduce aborts in OCC

Overview of ROCOCO

1. Two-phase protocol

- Most pieces are executed at the second phase

2. Decentralized dependency tracking

- Servers track pieces' arrival order
- Identify non-serializable orders
- Deterministically reorder pieces

3. Offline workload checking

- Identifies safe workloads (common)
- Identifies small parts that need traditional approaches (rare)

Conclusion of ROCOCO

Traditional protocols perform poorly w/ contention

- OCC aborts & 2PL blocks

Rococo defers execution to enable reordering

- Strict serializability w/o aborting or blocking for common workloads

Rococo outperforms 2PL & OCC

- With growing contention
- Scales out