

15. Hardware Security

Hardware Features Designed for Security

SMEP/SMAP: Supervisor Mode Execution Prevention/Supervisor Mode Access

Prevention

Return-to-user Attack: 内核代码中某些函数指针被设为NULL，dereference时跳转到用户代码。通过SMEP/SMAP禁止内核访问/执行用户代码。

- SMEP: Allows pages to be protected from supervisor-mode instruction fetches. If SMEP = 1, OS cannot fetch instructions from application
- SMAP: Allows pages to be protected from supervisor-mode data accesses. If SMAP = 1, OS cannot access data at linear addresses of application
- ARM的类似技术: PAN, Privileged Access Never; PXN, Privileged eXecute Never; UAO, User Access Only

应用: 使用SMAP做进程内隔离, 将用户代码放在内核态, 需要保护的数据放在用户态

ret2dir Attack: 每一页物理内存都同时拥有kernel和user两个虚拟地址, 虽然禁止内核执行用户代码, 但通过跳转到恶意代码对应的kernel地址仍可攻击。

MPX: Memory Protection eXtension

Background: C/C++ bounds error (可通过gcc的-fcheck-pointer-bounds flag防止)

Intel introduces MPX since Skylake

- Specified by two 64-bit addresses specifying the beginning and the end of a range
- New instructions are introduced to efficiently compare a given value against the bounds, raising an exception when the value does not fall within the permitted range
 - bndmov: Fetch the bounds information (upper and lower) out of memory and put it in a bounds register.
 - bndcl: Check the lower bounds against an argument (%rax)
 - bndcu: Check the upper bounds against an argument (%rax)
 - bnd retq: Not a "true" Intel MPX instruction
- usage: make CFLAGS="-mmpx -fcheck-pointer-bounds -lmpx" LDFLAGS="-lmpxwrappers -lmpx"
- 将bounds存在bounds table(a two-level radix tree)中, 最坏情况可多消耗400%内存, 且使用较多bound时性能降低

MPK: Memory Protection Keys

- With MPK, every page belongs to one of 16 domains. A domain is determined by 4 bits in every page-table entry (referred to as the protection key)
- For every domain, there are two bits in a special register (pkru), which denote whether pages associated with that key can be read or written
- Only the kernel can change the key of a page, application can read and write the pkru register using the rdpkru and wrpkru instructions respectively
 - use case 1: protect critical data within one address space
 - use case 2: prevent data corruption (In-memory database prevents writes most of the time, only enable changing data when needs to change)

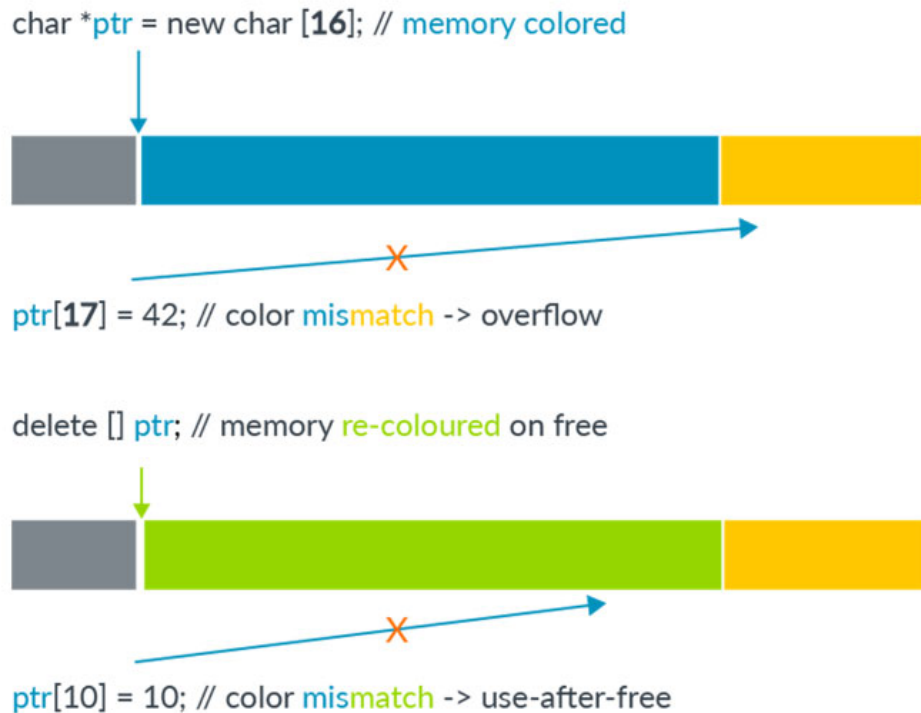
ARM PA: Pointer Authentication

- ARM只使用了64位中的40位, 可以用key对这40位地址进行加密, 保存在前24位, 从而作为验证, 防止地址被篡改
- PA defines five keys
 - Four keys for PAC* and AUT* instructions (combination of instruction/data and A/B keys),
 - One key for use with the general purpose PACGA instruction

- Keys are stored in internal registers and are not accessible by EL0 (user mode)
- New instructions:
 - PAC value creation: Write the value to the uppermost bits in a destination register alongside an address pointer value
 - Authentication: Validate a PAC and update the destination register with a correct or corrupt address pointer. If the authentication fails, an indirect branch or load that uses the authenticated, and corrupt, address will cause an exception

ARM MTE: Memory Tag Extension

Spatial Safety & Temporal Safety



- 引入新的内存类型: Normal Tagged Memory, 只允许指针访问相同tag的内存
- tag占4位, 即0~15
- 不是所有内存访问都需要tag checking, 如instruction fetches, translation table walks等
- MTE and PA可结合使用

Intel CET: Control-flow Enforcement Technology

- Shadow stack: a second stack for the program
 - 该栈只记录控制数据(return address), 与原有的stack同时push和pop, 如果RET时两个地址不一样, 则引发 control protection exception
 - 被页表保护
- Indirect Branch Tracking: New instruction: ENDBRANCH
 - call/jmp跳转到的指令必须以ENDBRANCH开头, 否则无效
 - cpu维护了一个状态机跟踪indirect call/jmp, 当出现这些指令, 状态由IDLE转为WAIT_FOR_ENDBRANCH, 该状态下一条指令必须为ENDBRANCH, 否则报错

Trusted Execution Environment

XOM: eXecute-Only Memory

- 代码和数据在内存中加密
- 存储加密值的哈希

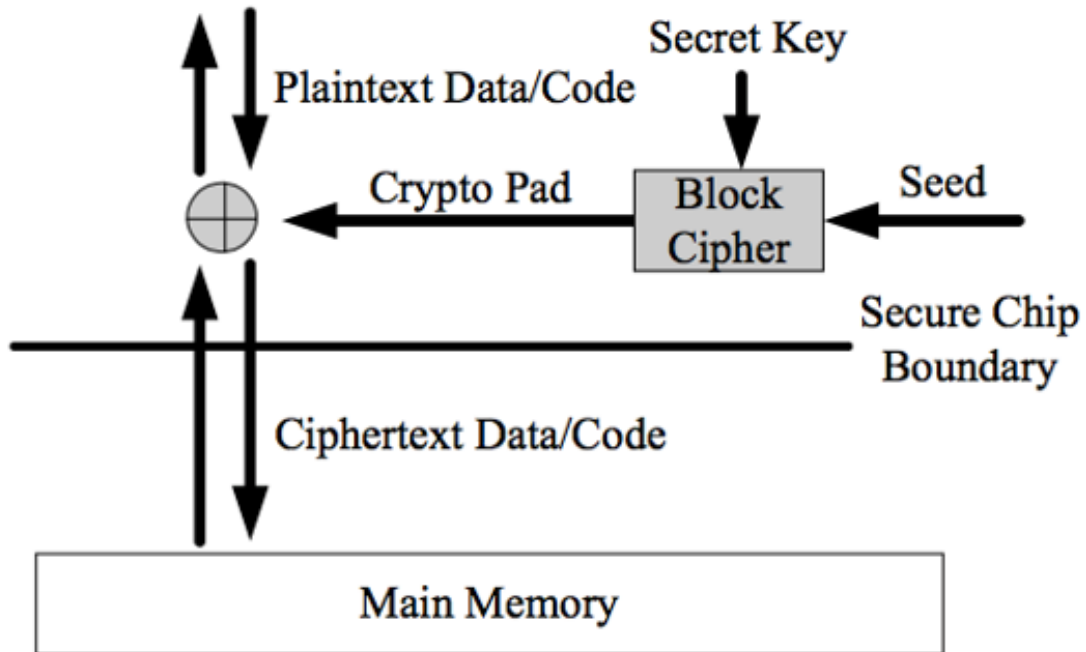


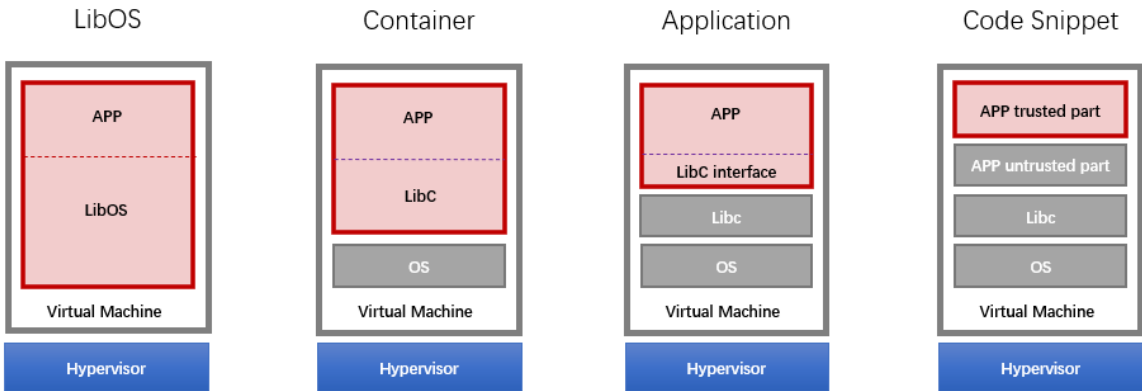
Figure 1. Counter-mode based memory encryption.

.

Intel SGX

- SGX Execution Flow
 - App built with trusted and untrusted parts
 - App runs & creates the enclave which is placed in trusted memory
 - Trusted function is called, execution transitioned to the enclave
 - Enclave sees all process data in clear; external access to enclave data is denied
 - Trusted function returns; enclave data remains in trusted memory
 - Application continues normal execution

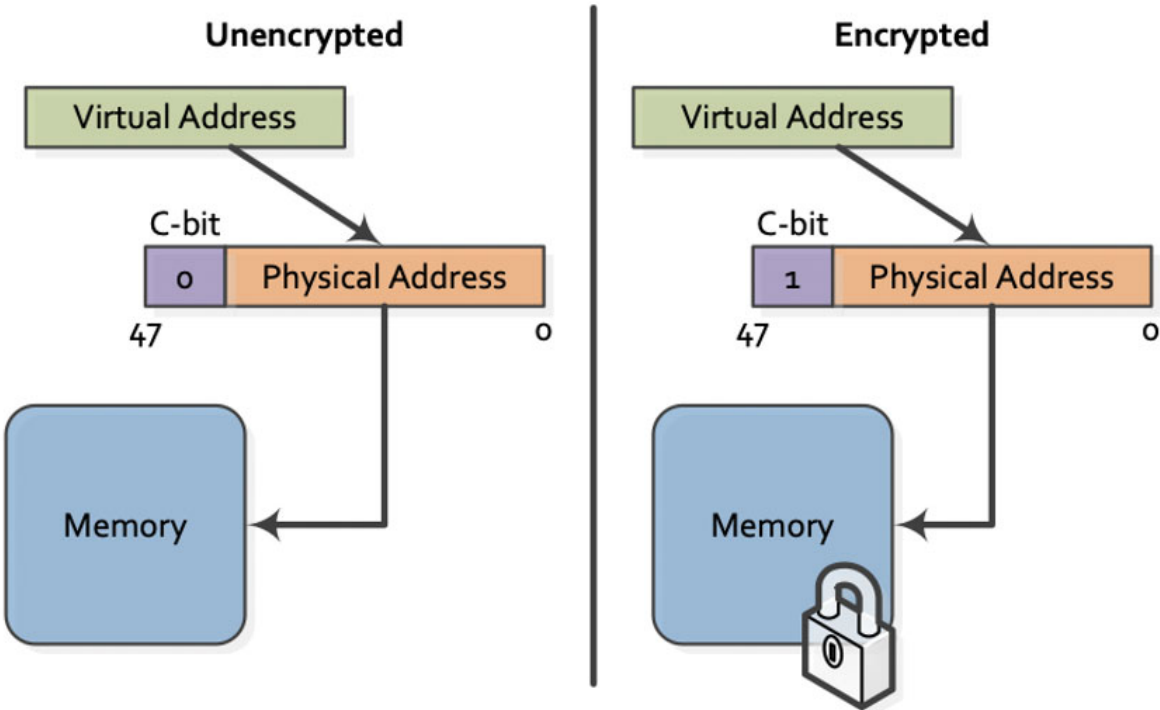
Software Architectures of SGX



	Compatibility	TCB Size	Ocall Num	Attack Surface	Protect OS
LibOS	Part	Large	Few	Small	✓
Container	✓	Mid	Mid	Mid	✗
Application	✓	Mid	Many	Large	✗
Code Snippet	✗	Small	Few	Small	✗

AMD's SME/SEV: Secure Memory Encryption/Secure Encrypted Virtualization

- Features



-
- Hardware AES engine located in the memory controller performs inline encryption and decryption of DRAM
- Minimal performance impact: Extra latency only taken for encrypted pages
- No application changes required
- Encryption keys are managed by the AMD Secure Processor and are hardware isolated, not known to any software on the CPU
- Comparing with Intel SGX
 - SME不会防范内核，用来防御cold-boot attacks，以及非易失内存的数据泄露

- SEV专注于虚拟机，可以防御其他虚拟机以及宿主机
- 类似技术: Intel MKTME: Multi-Key Total Memory Encryption, 支持多个key加密

ARM TrustZone

- Two modes: Normal world (REE, rich execution environment) and secure world(TEE, trusted execution environment), SMC instruction to switch
- 总线上增加1位，外设可以区分请求来自哪个world
- 应用: 手机指纹识别、交通工具、无人机禁飞区

RISC-V PMP/sPMP: Physical Memory Protection

- 通过一组PMP registers将物理内存划分为互相隔离的区域，每一段属于一个enclave
- 由于PMP registers数量有限，enclave数量也受限

Penglai

- Enclave on RISC-V ISA
- 为了实现细粒度内存划分，在DRAM增加一个bitmap，每一位表示一个页是否安全
- 所有不安全的页存储在一个隔离的内存区域PT_AREA
- 使用cache partition防御侧信道攻击
- 通过签名证明确实运行在enclave

Hardware Features Not Designed for Security

Intel TSX: Transactional Synchronization eXtensions

- Programming with RTM(restricted transactional memory)
 - If transaction starts successfully, do work protected by RTM, and then try to commit
 - If abort, system rollback to _xbegin, return an abort code
 - Manually abort inside a transaction
- 使用HTM保护数据: 将数据放在transaction中，利用HTM保证的原子性防止其他并发访问
- 利用HTM攻击KASLR: KASLR技术用来随机化内核地址。用户态随机访问一个地址，有两种情况，一是未映射，二是内核地址空间，两者返回segmentation fault有时间差。将这种试探代码放在transaction中，abort速度极快，可以快速试探出内核地址

Intel CAT: Cache Allocation Technology

The "Noisy Neighbor" Problem: 某些应用使用大量内存，但本身对cache需求不高（例如流媒体播放，之前的帧没有用处），反而占据了其他需要利用cache的应用的cache容量，影响性能。

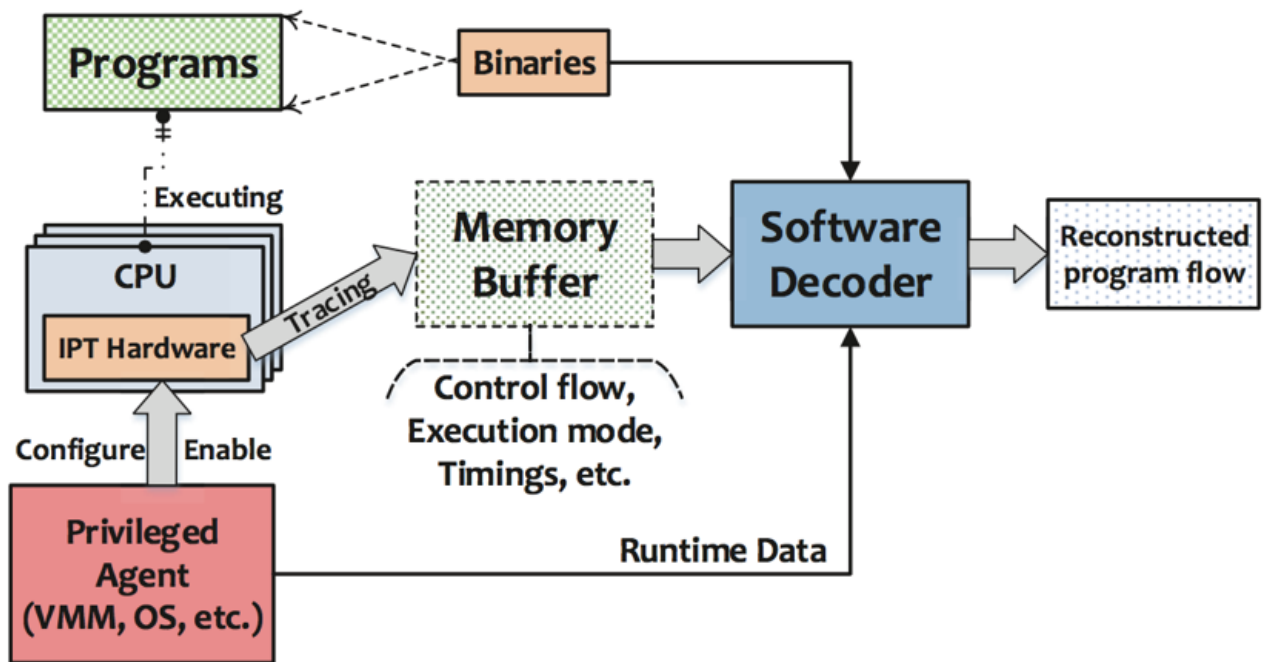
- CAT将thread / app / VM / container通过Class of Service (CLOS)分组，每个CLOS有相应的resource capacity bitmasks (CBMs)，表示可以使用cache的哪一部分
- 通过CAT防御基于cache的侧信道攻击（The PRIME+PROBE Attack，先用随机数据将cache填满，然后触发加密程序，最后重新访问原数据，通过cache miss情况推算加密行为），将不可信应用隔离，只能使用部分cache

PMU: Performance Monitor Unit

- BTS: Branch Trace Store, 记录程序所有地址跳转
- Motivation: Code Injection Attack, Code Reuse Attack
- 做法: 利用PMU监测CFI
 - Offline phase: 记录所有可能的分支跳转
 - Online phase: 将实际跳转与合法跳转对比，发现恶意行为
 - 记录3种合法跳转
 - ret_set: all the addresses next to a call
 - call_set: all the first addresses of a function

- train_sets: all the target addresses that once happened

Intel PT: Intel Processor Tracing



- 增加了硬件对Trace进行压缩，使tracing很快，但decode很慢
- FlowGuard: 将压缩的实际trace与压缩的可信trace直接对比，如果无法判断再解压缩

16. Data Privacy

ZKP: Zero-Knowledge Proof

Problem: Alice想向Bob证明她has the answer A of the problem P, 如果直接发送A, 则Bob也知道了A

Zero-Knowledge Proof:

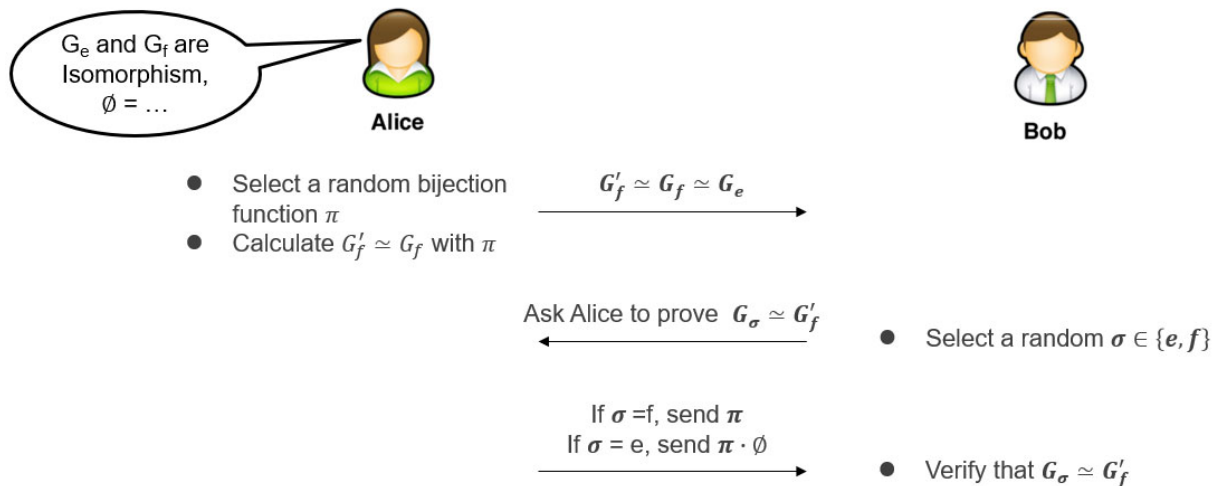
- Completeness: Alice can construct the proof if she has A
- Soundness: Alice cannot construct the proof if she doesn't have A
- Zeroknowledge: Bob knows nothing about A

Interactive Zero-Knowledge Proof: P has answer x of a problem L , and tries to prove it with > 1 iterations:

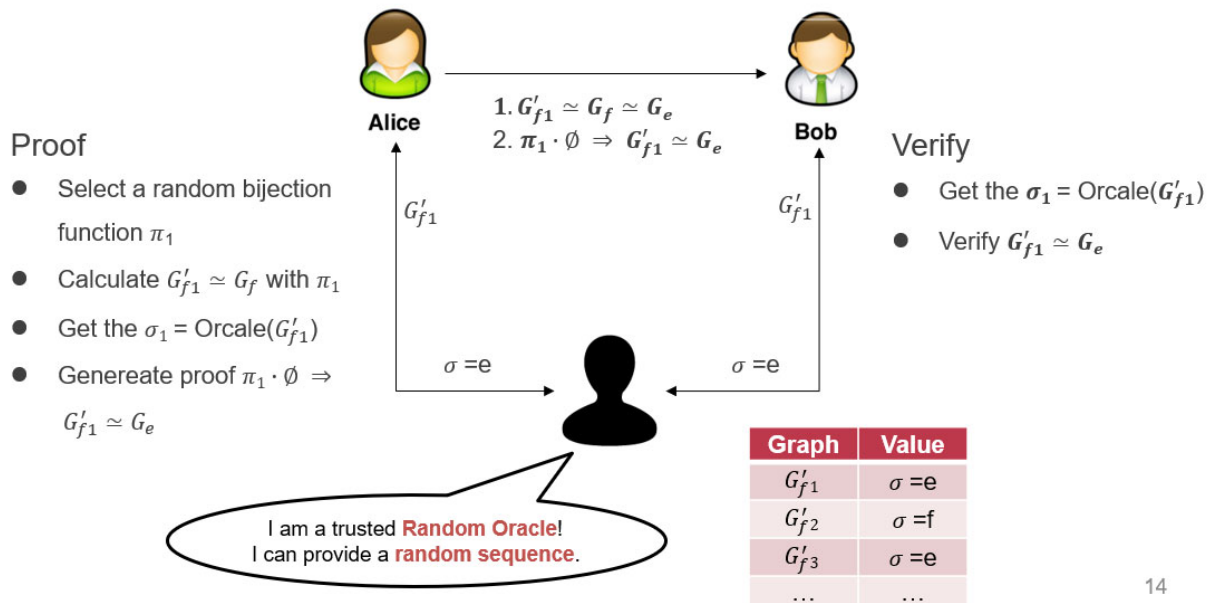
- Step-1: P transfers L to L' , and promises that L' is transferred from L and she has the answer x'
- Step-2: V challenges P
- Step-3: P shows the proof of the answer x' , which will not leak x
- V trusts that P has x when P always meets the challenge

图的同构: If $G_1 = (V_1, E_1)$ and $G_2 = (V_2, E_2)$ are isomorphic, there exist a bijection function (双射, 即——对应函数) ϕ , that for any $(u, v) \in E_1$, exist $\phi(u, v) \in E_2$, 即两张图的每条边都——对应

- 基于图同构的Interactive ZKP
 - Alice生成一个与两张图都同构的新图, 反复询问迭代



- Non-Interactive ZKP
 - 有一个可信的第三方生成随机序列, 让Alice一次性证明

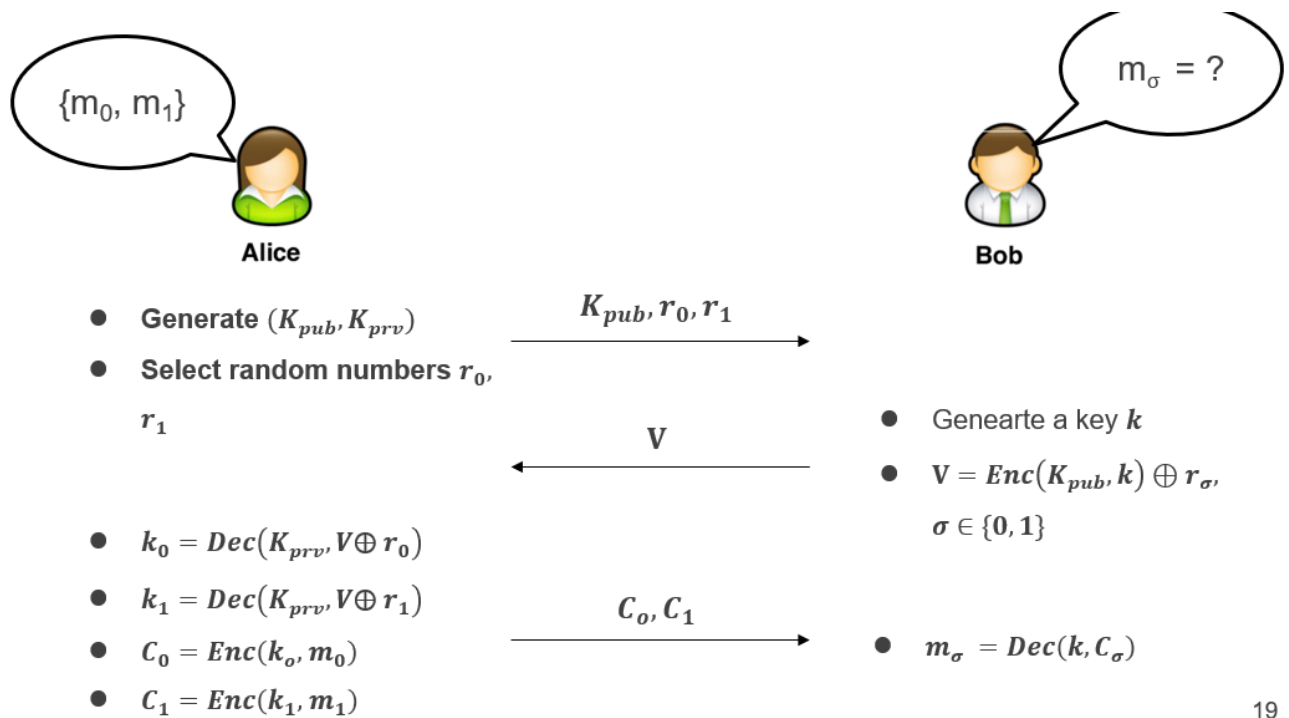


14

OP: Oblivious Transfer 不经意传输

Problem: sender不想让receiver拿到所有数据, receiver不想让sender知道自己想要哪个数据

1-out-of-2 OT:



19

- Bob把自己的key用Alice的公钥加密, 并与 r_i 作异或
- Alice分别用两个解密结果加密她的两份消息, 将加密结果送给Bob (只有Bob选择的 i 才能解密出Bob的key, 且Alice不知道Bob选择了哪个)
- Bob用 k 解密两份密文, 得到需要的消息, 而另一份无法解密

HE: Homomorphic Encryption 同态加密

Problem: 想把数据放在云端计算, 但想加密

Solution: 密文可以直接运算, 返回结果后用户解密即可。但目前算法难以支持所有密文运算, 例如RSA只支持乘法

SMPC: Secure Multi-Party Computing

Problem: 多方拥有不同数据, 想用这些数据共同计算而又不将数据泄露给其他方

Yao's Protocol: GC(Garbled Circuits)+OT(Oblivious Transfer), 姚期智的混淆电路+不经意传输算法

TEE: Trusted Execution Environment

- Software TEE
 - VM-based TEE
 - Same privilege protection
- ARM TrustZone
- Intel SGX
- AMD SME/SEV
- Penglai, SANCTUM

DP: Differential privacy 差分隐私

Problem: 想要拒绝用户访问数据库单个条目, 但不能简单限制, 否则可以通过 $\text{sum}(\cdot) - \text{sum}(\text{where } \neq \text{xx})$ 反推出

Solution: (若两个数据集有且仅有一条数据不一样, 则称此二者为相邻数据集) 如果某算法作用于任何相邻数据集, 得到一个特定输出的概率应差不多, 那么我们就说这个算法能达到差分隐私的效果。也就是说, 观察者通过观察输出结果很难察觉出数据集一点微小的变化, 从而达到保护隐私的目的。差分度越低, 安全性越高。

- 实现方式: 向计算函数中加噪声

FL: Federated Learning 联邦学习

Problem: 多方训练同一个模型, 需要保证各方数据隐私

- 横向: 一方拥有一个样本的全部数据
 - 服务器分发子模型给每个用户, 用户提交update, 会泄露个人隐私
 - 用户向update中添加噪声, 并保证总体对称, 相互抵消, 不影响总模型更新
 - 如果有用户掉线, 服务器必须向其他用户询问该用户的噪声偏移量; 则服务器可以伪造用户掉线, 通过询问计算出该用户原本的update
 - Secret Sharing & Double Masking算法
 - Share a secret S among N nodes, T nodes can reconstruct the secret, 基于k次函数方程求解需要k以上个点的原理
 - Each user u generates a_u and $S_{u,v}$ (for each user pair (u,v))
 - Each user u updates y_u , 是关于 a 和 S 的函数
 - Server计算 $\sum y_i$
 - For online node u , server asks other nodes to get the secret share of the a_u
 - For offline node u , server asks other nodes to get the secret share of the $S_{u,i}$
 - 所有节点防止server同时得到 a_u and $S_{u,i}$
- 纵向: 每一方拥有同一个样本的部分数据
 - 每方分别单独训练, 在全连接层汇总, 由有label的一方算出最终模型

后面介绍了3篇论文, 是对以上技术的具体应用

- Oblivious Multi-Party Machine Learning on Trusted Processors (Security'2016)
- BatchCrypt: Efficient Homomorphic Encryption for Cross-Silo Federated Learning (ATC'20)
- Privacy Accounting and Quality Control in the Sage Differentially Private ML Platform (SOSP'19)